

Network Security Class

Lab Session 2

Stefano Zanella - 621796

July 2013

1 Assignment

Experiment with the encoding, transmission and decoding in various realizations of the UEC, BSC, AWGN channels and for different parameters. Evaluate quantitatively and illustrate with plots and graphs:

1. the reliability (in terms of Bob's bit error rate on the secret message) and the secrecy (both in terms of Eve's bit error rate and information on the secret message) for the binary symmetric channel, over a wide range of bit error rates of the channels.
2. the reliability (in terms of Bob's bit error rate on the secret message) and the secrecy (both in terms of Eve's bit error rate and information on the secret message) for the additive white Gaussian noise, over a wide range of signal to noise ratios of the channels.

Discuss your results, in relationship with the channel secrecy capacity.

2 Uniform Error Channel

We start with the ideal case of a UEC. The given model sets a codeword length l_x of 7 bits; also it defines two different values for the number of errors that the legitimate and eavesdropper channels make during transmission, which are respectively $n_{err_B} = 1$, $n_{err_E} = 3$.

Theoretical analysis

From given parameters, we can derive the other variables that let us fully characterize the model, thus allowing for precise analysis of secrecy metrics. We start from calculating the value of $N_{y|x}$ and $N_{z|x}$, which are the cardinalities of the sets $T_{y|x}(a)$ and $T_{z|x}(a)$ that determine the possible messages Bob and Eve can receive given a determined $a \in \mathcal{X}$ and their respective number of possible errors. Using this definition, it's easy to retrieve the two values: that is, given n_{err_B} and a generic 7 bit codeword $a \in \mathcal{X}$, on the other side Bob can receive all 7 bit codewords containing 0 up to n_{err_B} errors. This value is represented as:

$$N_{y|x} = \sum_{i=0}^{n_{err_B}} \binom{7}{i} = 8$$

With the same approach we say that

$$N_{z|x} = \sum_{i=0}^{n_{err_E}} \binom{7}{i} = 64$$

That said, we can calculate the number of possible secret messages. We know from the theory of UECs and general physical layer secrecy that

$$|\mathcal{M}| \leq \frac{|\mathcal{Y}| / N_{y|x}}{|\mathcal{Z}| / N_{z|x}} = \frac{N_{z|x}}{N_{y|x}} = 8$$

which leads to

$$l_u = \log_2 8 = 3$$

That is, we can use 3 out of the 7 available bits to send secret information. This means we can achieve a secrecy rate of:

$$R_s = \frac{\log_2 |\mathcal{M}|}{n} = \frac{l_u}{l_x} = \frac{3}{7} \simeq 0.4286$$

Also, we can say that the channel has a secrecy capacity of:

$$\begin{aligned}
C_s &= I(x, y) - I(x, z) \\
&= (H(x) - H(x|y)) - (H(x) - H(x|z)) \\
&= H(x|z) - H(x|y)
\end{aligned}$$

which can be described, in the case of a UEC, as:

$$\begin{aligned}
C_s &= -\log_2 \frac{2^{l_x}}{N_{z|x}} + \log_2 \frac{2^{l_x}}{N_{y|x}} \\
&= \log_2 \frac{N_{z|x}}{2^{l_x}} - \log_2 \frac{N_{y|x}}{2^{l_x}} \\
&= \frac{\log_2 N_{z|x} - \log_2 N_{y|x}}{l_x} \\
&= \frac{\log_2 \frac{N_{z|x}}{N_{y|x}}}{l_x} \\
&= \frac{\log_2 8}{7} = \frac{3}{7} \simeq 0.4286
\end{aligned}$$

This means that in the case of an UEC we can aim at transmitting secret information filling the channel secrecy capacity.

Designing the encoder and decoder

At this point we've defined what are the achievable secrecy metrics in theory; we still need to actually design the encoder and the decoder so that we can do an actual simulation of secure transmission over the UEC.

We recap the basic features of (secure) physical layer encoder and decoder:

- The encoder is probabilistic, defined by $P_{x|u}(a|d)$ $a \in \mathcal{X}, d \in \mathcal{M}$
- The decoder is deterministic, defined by the function $\hat{u} = D(y)$ $D : \mathcal{Y} \rightarrow \mathcal{M}$
- The encoding/decoding process is **correct**: $P[\hat{u} \neq u] = 0$
- u is independent of the eavesdropped z (**secrecy** requirement).

We can divide the design process in two steps: coding for *reliability* (correctness) and coding for *secrecy*. In the first step we design the code \mathcal{C}' so

that it obeys to the property that each codeword $a \in \mathcal{X}'$ maps to a subset $T_{y|x}(a) \subset \mathcal{Y}$ and $\forall(a_1, a_2) \in \mathcal{X}', a_1 \neq a_2, T_{y|x}(a_1) \cap T_{y|x}(a_2) = \emptyset$. In the second step, instead, we refine the result of the preceding step so that the eavesdropper message space \mathcal{Z} is filled in such a way to make indistinguishable which the original message $d \in \mathcal{M}$ was (from an attacker's perspective).

To obtain the first property, since $n_{err_B} = 1$ we can rely on Hamming codes; for a given message a , its corresponding subset in \mathcal{Y} must contain it plus all messages with Hamming distance at most n_{err_B} from it. Also, we need to make every subset to be at distance $> 2n_{err_B}$ to achieve correctness. Since we are talking about binary codes, this translates into saying that the minimum distance between subsets is $2n_{err_B} + 1 = 3$. A code that satisfies this property is **Hamming(7,4)**; it allows not only to define disjoint subsets over \mathcal{Y} , but to actually define a *partition* of \mathcal{Y} .

That said, we can now consider code secrecy with respect to the eavesdropper. As already said, such a code should fill the eavesdropper message space \mathcal{Z} . We already saw that $N_{z|x} = 64$; this means that the Z set is partitioned into $2^{l_x}/N_{z|x} = 2$ subsets.

A way to define these two subsets is to consider messages which are one the complement of the other and put them in different subsets. That is, messages at $d_H = 7$ belong to different subsets. It's easy to see how this is an equipartition of \mathcal{Z} , thus obeying to the properties of the UEC.

That said, we need a way to design the encoder so that for every message $d \in \mathcal{M}$ there is an equal probability that its encoding belongs to each of the two subsets. This can be done as follows:

- pick a generic message $u = u_1 \ u_2 \ u_3$ and choose a bit b at random: $b \sim \mathcal{U}(\mathbf{B})$.

- let

$$v = \begin{cases} b|u & \text{if } b = 0 \\ b|\bar{u} & \text{if } b = 1 \end{cases}$$

- encode v with **Hamming(7,4)**

It turns out that the two choices for v are complementary, and so are the **Hamming(7,4)** codewords for them. This means that:

- the encoder is probabilistic as requested:

$$P_{x|u}(a|d) = \begin{cases} 1/2 & \text{if } a \in T_{x|u}(d) \\ 0 & \text{otherwise} \end{cases}$$

- the decoder is deterministic: since it has $n_{err_B} = 1$, the `Hamming(7,4)` code allows for correct recovery of the original message x . Then, the decoder can just discard the 3 bits prepended by the Hamming encoder and retrieve v . At this point, it can look at the first bit b and decode:

$$\hat{u} = \begin{cases} v_2 v_3 v_4 & \text{if } b = 0 \\ \overline{v_2 v_3 v_4} & \text{if } b = 1 \end{cases}$$

- eavesdropper message space is covered by each and every message $d \in \mathcal{M}$; furthermore, each of its subsets is disjoint from each other.
- legitimate receiver message space is divided into disjoint subsets.

So, mathematically, the transmission with the designed encoder is both *reliable* and *secure*. We now need to actually measure secrecy and reliability with the Matlab simulation.

Transmission simulation

The simulation is performed by `UECscript.m`. It accepts two parameters: a number of trials and a number representing how many times to repeat that number of trials in a cumulative way. It basically simulates `trials · split` transmissions over the UEC, counting the number of errors made by the legitimate receiver and the eavesdropper. Then, every `trials` transmissions, it calculates the BER for the two transmission ends and print an histogram representing residual information on u given a particular z for every codeword x , as well as information about messages, joint and mutual entropies.

We start by observing that we need to perform a certain number of transmissions before reaching the bounds given by the statistical model. In particular, this can be noticed when looking at the measured conditional entropy between u and z . Looking at Figure 1 we see that we approach asymptotically the theoretical result of statistical independence between u and z (which means $H_{u|z} = l_u = 3$); to obtain a good approximation, we need to consider at least 8-10000 samples (to have $\delta \sim 2\%$).

The same result can be observed comparing the histograms depicting conditional entropy for every codeword fixed $z = 0$ with different number of trials. That is, if we compare Figure 2, 3 and 4, we see that with only 1000 experiments there are certain codewords that let leak 2 bits of the original




Figure 1: $H_{u|z}$ vs. number of trials

message. This is clearly not a flaw of the model, but it rather points an insufficient amount of data; this is confirmed by looking at what happens with 5000 and 10000 experiments; we see that the distribution among the codewords tends to become flat and within 0.25 bit from the theoretical result.

At the same time, we can see another interesting result. Conversely from the case of conditional entropy, if we plot the BER against the number of experiments for both the legitimate receiver and the eavesdropper as in Figure 5, we see that no matter how many experiments we do:

- The BER at the legitimate receiver is always 0
- The BER at the eavesdropper fluctuates tightly around 0.5




Figure 2: $H_{u|z=0}$ with 1000 trials

This is an expected result and proves the design of the encoder is done right. In fact, the first observation tells us that the codec achieves *correctness*. This is expected since we used an 1-bit error-correction code and the legitimate receiver can make *at most* 1 error.

Instead, the second observation tells us that we achieve the goal of confusing the attacker as much as possible: in fact, achieving a BER of 0.5 means that, for every received bit, the attacker has no clue if it has been corrupted during transmission or not, and so he has no better choice than randomly guessing the eavesdropped secret message.

Also, a note on BER and conditional entropy: we cannot claim we've achieved (perfect) secrecy just by looking at the BER (it's a necessary condition, but not sufficient); this is why we also take into consideration measured




Figure 3: $H_{u|z=0}$ with 5000 trials

conditional entropy between eavesdropped and original messages. This is because BER is just a global measure over channel’s physical behaviour, but doesn’t take into account the *semantic* of the error, that is, it doesn’t provide any information over actual *information*. We can partially see this in the case of a low number of experiments: while we can achieve $\text{BER} = 0.5$, we cannot achieve a good conditional entropy.

Instead, the notion of perfect secrecy states explicitly that original and eavesdropped messages must be statistically independent: that is, knowing the latter doesn’t give any advantage over knowledge of the former. By calculating and plotting conditional entropy we’re actually measuring this.

Below, we recap the result of the simulation of 10000 transmissions, as output by the `UECscript` function:




Figure 4: $H_{u|z=0}$ with 10000 trials

Secrecy capacity = 0.4286
 Secrecy rate = 0.4286
 BER at Bob = 0.00e+00
 BER at Eve = 5.03e-01

 H(u) = 2.9996 bit
 H(z) = 6.9889 bit
 H(u,z) = 9.9256 bit
 H(u|z) = 2.9367 bit
 H(z|u) = 6.9260 bit
 I(u;z) = 0.0629 bit

This result will serve as a reference for the models that follow.




Figure 5: BER for E and B over a UEC

3 Binary Symmetric Channel

We now look at what happens when we model the channel as a BSC, and see how secrecy parameters change when error probabilities change.

We use the same encoding scheme as before, so everything said about it still holds.

Secrecy Capacity

To start, we try to model a channel as close as possible to the UEC, in terms of introduced noise. Specifically, we can set the error probabilities to:¹

$$\varepsilon_B = n_{err_B}/l_x = 1/7 \simeq 0.14286$$

$$\varepsilon_E = n_{err_E}/l_x = 3/7 \simeq 0.42857$$

In this case, it's easy to see that:

$$C_s = h_2(\varepsilon_E) - h_2(\varepsilon_B) = 0.39356$$

which is lower than what we achieved with the UEC and corresponds to:

$$C_s \cdot l_x = 2.7549 \text{ bit}$$

that can be concealed in a single message. To achieve a secrecy capacity that allows to conceal a 3 bit message, we can try to lower ε_B , increase ε_E , or both. We see that even if we increase ε_E so to achieve maximum confusion, i.e. $\varepsilon_E = 0.5$, we can only reach $C_s = 2.8583$. Instead, we can reach our goal by leaving $\varepsilon_E = 0.42857$ and setting $\varepsilon_B = 0.125$, which lead to $C_s \cdot l_x = 3.0916$ bit. In this case, the simulation of 10000 transmissions leads to the following result:

```
Metrics with 10000 trials:
Secrecy capacity = 0.4417
Secrecy rate      = 0.4286
BER at Bob       = 1.26e-01
BER at Eve       = 4.88e-01
-----
H(u)    = 2.9989 bit
H(z)    = 6.9923 bit
H(u,z)  = 9.9222 bit
H(u|z)  = 2.9300 bit
H(z|u)  = 6.9233 bit
I(u;z)  = 0.0690 bit
```

Also, in Figure 11 it can be seen how conditional entropy is distributed, as in the case of the UEC.

¹This doesn't mean there's a strict equivalence between the values of ε_B and ε_E , and the error rates of the UEC. Actually, in the case of the UEC the error rate is considered on the *whole message*, while the error probability of the BSC is applied *independently for every bit sent*. Though, this simple relation helps targeting a starting point for the analysis.




Figure 6: $H_{u|z=0}$ with 10000 trials

We can make two key observations here:

- the model doesn't achieve correctness ($\text{BER}_B > 0$)
- albeit the level of secrecy is high both in terms of BER and statistical independence of u and z , the error probability for eavesdropper's channel is very high, meaning the channel is almost unusable ($0.5 =$ highest level of confusion).

The first one is an important difference from UEC: since every bit transmitted over a BSC has the same probability ε of being swapped, we can't be sure, as in the case of UEC, that the legitimate channel makes at most a fixed amount of errors. In the particular case of our `Hamming(7,4)` code, we can calculate the probability of correctly receiving a message as the probability

of receiving a message with at most 1 error:

$$p[\hat{u} = u]_B = (1 - \varepsilon_B)^7 + (1 - \varepsilon_B)^6 \cdot \varepsilon_B \simeq 0.44880$$

which turns out to be less than %50 percent of the times.

On the contrary, on the eavesdropper side, the probability of receiving an unrecoverable message is:

$$\begin{aligned} p[\hat{u} \neq u]_E &= 1 - p[\hat{u} = u]_E \\ &= 1 - [(1 - \varepsilon_E)^7 + (1 - \varepsilon_E)^6 \cdot \varepsilon_E] \\ &\simeq 0.99503 \end{aligned}$$

which is extremely high, as already pointed out.

We can see the trend of these two functions (isomorph to each other) in Figure 7.

From that, we can also try to see how secrecy capacity evolves when varying both ε_B and ε_E . Figure 8 shows the 3D curve representing C_s as a function of both variable (when $\varepsilon_B < \varepsilon_E$, since in the other case the eavesdropper would have an advantage over the legitimate receiver and the roles could be considered swapped).

From Figure 8 we can confirm what intuition already pointed out:

- we can't have any form of secrecy when the error probabilities of the two channel are equal;
- full channel secrecy can be reached only in the (rather theoretical) case of $\varepsilon_B = 0, \varepsilon_E = 0.5$;
- as error probability for legitimate channel increases, no matter the error probability for the eavesdropper, we can only hope to obtain a certain amount of secrecy that decreases with the former.

That said, we can limit the observation to the cases that give us a secrecy capacity high enough to conceal the 3 bits we want to send ($C_s > 3/7$). Figure 9 shows the region that satisfies this condition.

From this plot, we can select few points and analyze them more in depth. To choose them, let's take for granted that we want to send 3 secure bits, and focus on the following criteria:

- points that lead to $R_s = C_s - \varepsilon$
- points that lead to $R_s = C_s + \varepsilon$

In particular:




Figure 7: Decoding probabilities for the **Hamming(7,4)** code vs ε

- $\varepsilon_B = 0.06, \varepsilon_E = 0.22$ ($C_s = 0.43272$)
- $\varepsilon_B = 0.12, \varepsilon_E = 0.38$ ($C_s = 0.42868$)
- $\varepsilon_B = 0.10, \varepsilon_E = 0.30$ ($C_s = 0.41230$)
- $\varepsilon_B = 0.06, \varepsilon_E = 0.22$ ($C_s = 0.31957$)




Figure 8: Evolution of secrecy capacity as a function of ε_B and ε_E




Figure 9: Evolution of C_s under the constraint $C_s > 3/7$

Number of experiments

Before starting the analysis, let's first determine how many transmission to do per instance of the model.

If we plot the BER and the conditional entropies vs. the number of trials for a fixed value of ε_B and ε_E (i.e. 0.1 and 0.35 respectively), we see (Figure 10 and 10) that:

- as in the previous case, the value of BER for both Bob and Eve doesn't have direct relationship with the number of trials
- the value of conditional entropy still has a relationship with the number of trials, but there is somewhat less variation in the range [5000, 10000] than in the case of UEC.

For these reasons, given that running a simulation with 10000 transmissions is quite expensive in terms of time, let's just reduce them to 6000.




Figure 10: BER for E and B over a BSC




Figure 11: $H_{u|z}$ vs. number of trials

$$R_s < C_s: \varepsilon_B = 0.06, \varepsilon_E = 0.22$$

The results of the simulation are summarized in Figure 12 and in the following output:

```
Metrics with 6000 trials:
Secrecy capacity = 0.4327
Secrecy rate      = 0.4286
BER at Bob        = 3.58e-02
BER at Eve        = 2.75e-01
-----
H(u)   = 2.9994 bit
H(z)   = 6.9468 bit
H(u,z) = 9.0126 bit
H(u|z) = 2.0657 bit
H(z|u) = 6.0131 bit
I(u;z) = 0.9337 bit
```

It can be seen that, despite transmitting at a rate lower than the overall channel capacity, almost a whole bit leaks on average on the eavesdropper side. Also the BER is not adequate on the eavesdropper side, as an error occurs only in the 27.5% of the cases. Furthermore, if we look at the plot, we see that there are many codewords that reveal more than 2 bits to the eavesdropper; clearly this result is not acceptable in terms of confidentiality.

Instead, reliability on the legitimate side, while not perfect, it's quite good as we get an error in less than 4% of the cases.

$$R_s < C_s: \varepsilon_B = 0.12, \varepsilon_E = 0.38$$

Running the script gives the following result:

```
Metrics with 6000 trials:
Secrecy capacity = 0.4287
Secrecy rate      = 0.4286
BER at Bob        = 1.12e-01
BER at Eve        = 4.63e-01
-----
H(u)   = 2.9989 bit
H(z)   = 6.9833 bit
H(u,z) = 9.8317 bit
H(u|z) = 2.8484 bit
```




Figure 12: $H_{u|z=0}$ for $\varepsilon_B = 0.06, \varepsilon_E = 0.22$

$$H(z|u) = 6.8328 \text{ bit}$$

$$I(u; z) = 0.1505 \text{ bit}$$

Despite the transmission almost fills the whole channel capacity (when in the previous case the secrecy rate was more lower than capacity) we achieve better results than the previous case. In particular, we almost decimated the amount of revealed information; we can also confuse the eavesdropper more, generating an error in more than 46% of the times. Also the plot in Figure 13 shows that all eavesdropped codewords don't retain more than 0.5 bit of mutual information with their original counterpart.

On the side of reliability, instead, things are much worse than before. In particular, it can be noticed how, doubling the error probability on Bob we




Figure 13: $H_{u|z=0}$ for $\varepsilon_B = 0.12, \varepsilon_E = 0.38$

have almost triplicated the error rate, which now is at 11%. Given this is the result to the net of all recoverable errors, it can be considered quite high to consider the transmission truly "reliable".

$$R_s > C_s: \varepsilon_B = 0.10, \varepsilon_E = 0.30$$

The results are:

Metrics with 6000 trials:

Secrecy capacity = 0.4123

Secrecy rate = 0.4286

BER at Bob = 8.74e-02

BER at Eve = 3.86e-01

 $H(u)$ = 2.9991 bit
 $H(z)$ = 6.9844 bit
 $H(u,z)$ = 9.6046 bit
 $H(u|z)$ = 2.6201 bit
 $H(z|u)$ = 6.6055 bit
 $I(u;z)$ = 0.3790 bit




Figure 14: $H_{u|z=0}$ for $\varepsilon_B = 0.10, \varepsilon_E = 0.30$

It is clear that we aren't achieving neither much reliability (even if better than the previous case), nor secrecy (Figure 14 shows that there are code-words that leak up to a whole bit to the eavesdropper). What is interesting, though, is the comparison of this outcome with that obtained in the first

case. In fact, despite transmitting at a secrecy rate higher than the secrecy capacity of the channel, despite having a ratio between error probabilities of 3 instead of almost 4 in the other case, we are *still* able to send more secure information than the case of $\varepsilon_B = 0.06, \varepsilon_E = 0.22$. This points out how important is for transmission confidentiality that the eavesdropper has a noisy channel.

$$R_s > C_s: \varepsilon_B = 0.08, \varepsilon_E = 0.20$$

The results are:

```
Metrics with 6000 trials:
Secrecy capacity = 0.3197
Secrecy rate      = 0.4286
BER at Bob        = 6.12e-02
BER at Eve        = 2.47e-01
-----
H(u)      = 2.9986 bit
H(z)      = 6.9236 bit
H(u,z)    = 8.8099 bit
H(u|z)    = 1.8863 bit
H(z|u)    = 5.8114 bit
I(u;z)    = 1.1123 bit
```

Here we are clearly in a bad condition. Probably the clearest evidence is Figure 15, which shows how there are codewords that are almost completely eavesdropped by Eve. As in the previous case, it's interesting to see how, despite transmitting at $R_s \gg C_s$, the confidentiality is not that much worse than the case of $\varepsilon_B = 0.06, \varepsilon_E = 0.22$, again pointing out the importance of the noise level on Eve's channel.

Further considerations

We can try to see what happens in cases in which we don't achieve high secrecy by reducing the number of secure bits transmitted. In particular, for $\varepsilon_B = 0.06, \varepsilon_E = 0.22$, if we set $l_u = 2$ we obtain:

```
Metrics with 6000 trials:
Secrecy capacity = 0.4327
Secrecy rate     = 0.2857
```





Figure 15: $H_{u|z=0}$ for $\varepsilon_B = 0.08, \varepsilon_E = 0.20$

BER at Bob = 3.62e-02
 BER at Eve = 2.77e-01

 $H(u)$ = 1.9997 bit
 $H(z)$ = 6.9349 bit
 $H(u,z)$ = 8.4239 bit
 $H(u|z)$ = 1.4890 bit
 $H(z|u)$ = 6.4242 bit
 $I(u;z)$ = 0.5108 bit

In the case $\varepsilon_B = 0.1, \varepsilon_E = 0.3, l_u = 2$, instead we obtain:

Metrics with 6000 trials:

```

Secrecy capacity = 0.4123
Secrecy rate     = 0.2857
BER at Bob      = 8.39e-02
BER at Eve      = 3.83e-01

```

```

-----

```

```

H(u)    = 1.9991 bit
H(z)    = 6.9819 bit
H(u,z)  = 8.7899 bit
H(u|z)  = 1.8079 bit
H(z|u)  = 6.7907 bit
I(u;z)  = 0.1912 bit

```

We see that we're able to obtain more statistical independence between u and z , in both cases (almost half of mutual information as the cases with $l_u = 3$). At the same time, though, it's clear how the BER isn't much different than before. This is evidence of what we've already stated when analyzing the UEC: BER isn't a sufficient condition to prove perfect secrecy.

Conclusions

Summing up all experimental observations, we can draw some conclusions before moving to the case of AWGN channels.

We have seen that, as we move toward the case $\varepsilon_B = 0, \varepsilon_E = 0.5$ we also make the channel behave as a UEC. Clearly, this is just a theoretical result, as those error probabilities are quite unlikely to find in practice. What happens far from these bounds is that all error patterns are possible on both channels, but they aren't distributed uniformly. This condition is then reflected in \mathcal{Y} and \mathcal{Z} , and their partitions $\{T_{y|x}\}$ and $\{T_{z|x}\}$; in particular, while both sets are actually partitioned such that to allow perfect secrecy, the distribution of the error patterns inside the subsets forming the partition is not uniform. It turns out that this condition is what avoid us to obtain perfect secrecy. In fact, the notion of perfect secrecy we've achieved in theory for the BSC relies on the notion of typical sequences, which provide only an asymptotical result, requiring very long transmissions in order to achieve uniformity. In our case, the combination of $l_x = 7$ and the number of simulated transmissions is not sufficient to guarantee long term, asymptotical uniformity.

Also, we've seen how in the case of BSCs, a fundamental condition is that of the eavesdropper's channel being noisy. If this condition is not met, no matter how high the secrecy capacity of the channel (i.e. how low the

error probability for the legitimate receiver), we're not able to conceal much information. This can be seen as another way to state that high BER on the eavesdropper side is a necessary condition for (perfect) secrecy.

4 AWGN Channel

Now we try to analyze the kind of secrecy achievable in a more realistic model. We consider the case of an AWGN channel, with different SNRs at the legitimate receiver end and at the eavesdropper end. We still consider the case of $l_x = 7, l_u = 4$ and still perform the same kind of encoding of the secret message u into the codeword x by means of the `Hamming(7,4)` code. Then, we take the codeword and perform a **128-PAM modulation** step. After that, we simulate transmission over an AWGN channel using Matlab/Octave builtin functions. Finally, we let both the eavesdropper and the legitimate receiver decode the received signal.

Clearly, in this case we are not considering the transmission over a binary medium anymore. Instead, we are dealing with a continuous signal; for this reason, we cannot apply the information theory functions and metrics (entropy, etc.) as we have done in the previous cases. So, the only metric we can rely upon for the analysis is BER, which, as we've already seen, it is just a necessary condition for (perfect) secrecy.

Number of transmissions

As we've already seen in both the previous cases, BER doesn't show much correlation with the number of transmissions over which we calculate it. So, for this case we limit the number of observations to 2000 per model instance. This makes each simulation faster, allowing for a broader range of trial values.

Secrecy Capacity

As we've done in the BSC case, we can derive a formulation of the secrecy capacity as a function of SNR_B and SNR_E and plot this function to find interesting regions to analyze. Though, in contrast with the other cases, here we can achieve (at least theoretically) unconstrained secrecy capacity.

This can be seen at first by looking at the expression for C_s :

$$C_s = \frac{1}{2}[\log_2(1 + \Lambda_B) - \log_2(1 + \Lambda_E)]$$

that is, if we let Λ_B increase and Λ_E decrease, the function will increase without bound.

For practical reasons, we'll limit our analysis within 0 and 50 dB.

That said, using the expression above for secrecy capacity, we can see how it evolves with respect to Λ_B and Λ_E in Figure 16.




Figure 16: Evolution of secrecy capacity as a function of Λ_B and Λ_E

As in the previous cases, here too we observe that secrecy capacity decreases as Λ_E increase, no matter how high the value for Λ_B is. We also see

that even in an extreme case like $\Lambda_B = 50$ dB, $\Lambda_E = 5$ dB we only reach $C_s \simeq 1.54$, while we send at

$$R_s = l_M \frac{l_u}{l_x} = 7 \cdot \frac{3}{7} = 3$$

Looking at this last expression, we can have a clue on a possible way to make the secrecy rate fit into channel's secrecy capacity. In fact, we see that in the expression it's included the ratio l_M/l_x , which is the code rate of the modulation system. In our case this ratio is 1, because we have a direct mapping between one message and the symbol used to perform amplitude modulation. But we can think of sending longer messages so that they're split into more symbols; in this case, we'll have $l_M/l_x < 1$, thus allowing for secrecy rate adjustment to fit available secrecy capacity. The problem, then, is that of finding a suitable combination of parameter $M = 2^{l_M}$ and codeword length l_x . For example, if we maintain the Hamming encoding as is, we need to use a length in the form $l_x = 2^r - 1$ (7, 15, 31, etc), which clearly doesn't allow for codewords to be split in an integer number of symbols. We can, however, think of modifying the encoding scheme (and decoding scheme accordingly) by padding the message with random bits so that $l'_x = k \cdot l_M$. For example, we can consider a 256-PAM modulation ($l_M = 8$) and the Hamming(15,11) encoding with a 1 bit pad so that $l'_x = 16$; this way we'll obtain $R_s = 8 \frac{3}{16} = 1.5$, fitting the secrecy capacity we already calculated. Note, however, that such technique incurs in a major drawback: by splitting the message into more symbols, we'll be able to send a complete message (and thus we'll be able to decode a secret message) once every l_M/l_x symbol periods. That is, by adjusting the code rate we incur in a tradeoff between secrecy rate and speed; and the situation is ever worse if we consider that we're not aiming at sending all l_x bits of information, but only the $l_u \ll l_x$ secret ones. In the example case (which is, again, very extreme) we'll end up with a 16 bit codeword to send a mere 3 bit secret message; so there's a 5+ times factor overhead.